

CSE 150B Week 3 Discussion

Chia-Han Chen, Raymond Xia, Benjamin Xia

Adversarial Search

Game Tree:

- Each node represents a state of the game
- Edges emanating from a node represent possible moves that a player can make at the state.
- Each player can make a finite number of actions, and the game has finite length.
- Assumes perfect information.

A very simple example: Tic Tac Toe

Let X be the max player


Let O be the min player


X's turn to play


Game score = 1 if X wins


Game score = -1 if O wins


Game score = 0 otherwise















Minimax values:

Node 2: +1 (terminal node)

Node 5: -1 (terminal node)

Node 7: 0 (terminal node)

Node 6: ?



Minimax values:


Node 2: +1 (terminal node)

Node 5: -1 (terminal node)

Node 7: 0 (terminal node)

Node 6: 0 (Max(Node 7))


Node 3: ?



Minimax values:

- Node 2: +1 (terminal node)
- Node 5: -1 (terminal node)
- Node 7: 0 (terminal node)
- Node 6: 0 (Max(Node 7))
- Node 3: -1 (Min(Node 5, 6))
- Node 1: ?

Hint: Minimax value of node 4 is 0



Minimax values:

Node 2: +1 (terminal node)

Node 5: -1 (terminal node)

Node 7: 0 (terminal node)


Node 6: 0 (Max(Node 7))


Node 3: -1 (Min(Node 5, 6))


Node 4: 0

Node 1: +1 (Max(Node 2, 3, 4))

| | | |
|--|--|--|
| | | |
| | | |
| | | |
| | | |








(Winner - X)

Minimax

Value of Nodes:



- Max Node: $\max(a_1, a_2, a_3, \dots, a_n)$



Minimax

Value of Nodes:




- Max Node: $\max(a_1, a_2, a_3, \dots, a_n)$
- Min Node: $\min(a_1, a_2, a_3, \dots, a_n)$



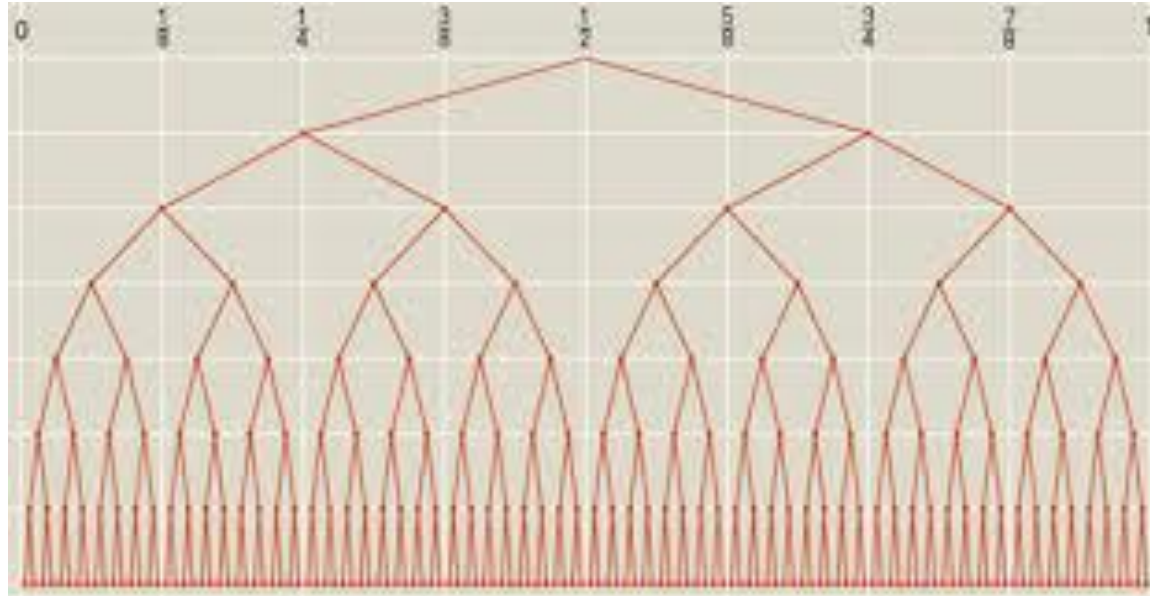
Minimax

Value of Nodes:

- Max Node: $\max(a_1, a_2, a_3, \dots, a_n)$
- Min Node: $\min(a_1, a_2, a_3, \dots, a_n)$
- Terminal Node (Leaf): Payoff/Outcome




Computing Minimax: Analysis



If each node has n possible actions, how many nodes do we have to explore to compute minimax for a tree with depth d ?

Recall: Depth is number of edges/actions from root node to terminal node

Computing Minimax: Analysis



If each node has n possible actions, how many nodes do we have to explore to compute minimax for a tree with depth d ?

$$1 + n + n^2 + n^3 + n^4 \dots = \sum_{i=0}^d n^i = \frac{n^{d+1}-1}{n-1}$$

In practice: Use alpha-beta pruning (Not applicable in expectimax) or heuristic after a certain depth.


Recall: Depth is number of edges/actions from root node to terminal node

Pruning

- Horizontal: Alpha-Beta Pruning
- Vertical: Cut-off at some fixed depth and use heuristics to evaluate the intermediate nodes.


Alpha – Beta Pruning

Alpha-Beta Pruning



Alpha – Beta Pruning

Alpha-Beta Pruning



Alpha – Beta Pruning

```
def alphabeta_minimax(node, alpha, beta):
    if terminal(node):
        return payoff(node)
    elif max_player(node):
        value = -infinity
        for n in children(node):
            value = max(value, alphabeta_minimax(n, alpha, beta))
            alpha = max(alpha, value) '''Try to push up'''
            if alpha >= beta: '''If alpha seems too big, stop'''
                break
        return value
    elif min_player(node):
        value = infinity
        for n in children(node):
            value = min(value, alphabeta_minimax(n, alpha, beta))
            beta = min(beta, value) '''Try to push down'''
            if beta <= alpha: '''If beta looks too small, stop'''
                break
        return value
```

First call:

```
alpha = -infinity '''fallback for Max'''
beta = infinity '''fallback for Min'''
node = root
```

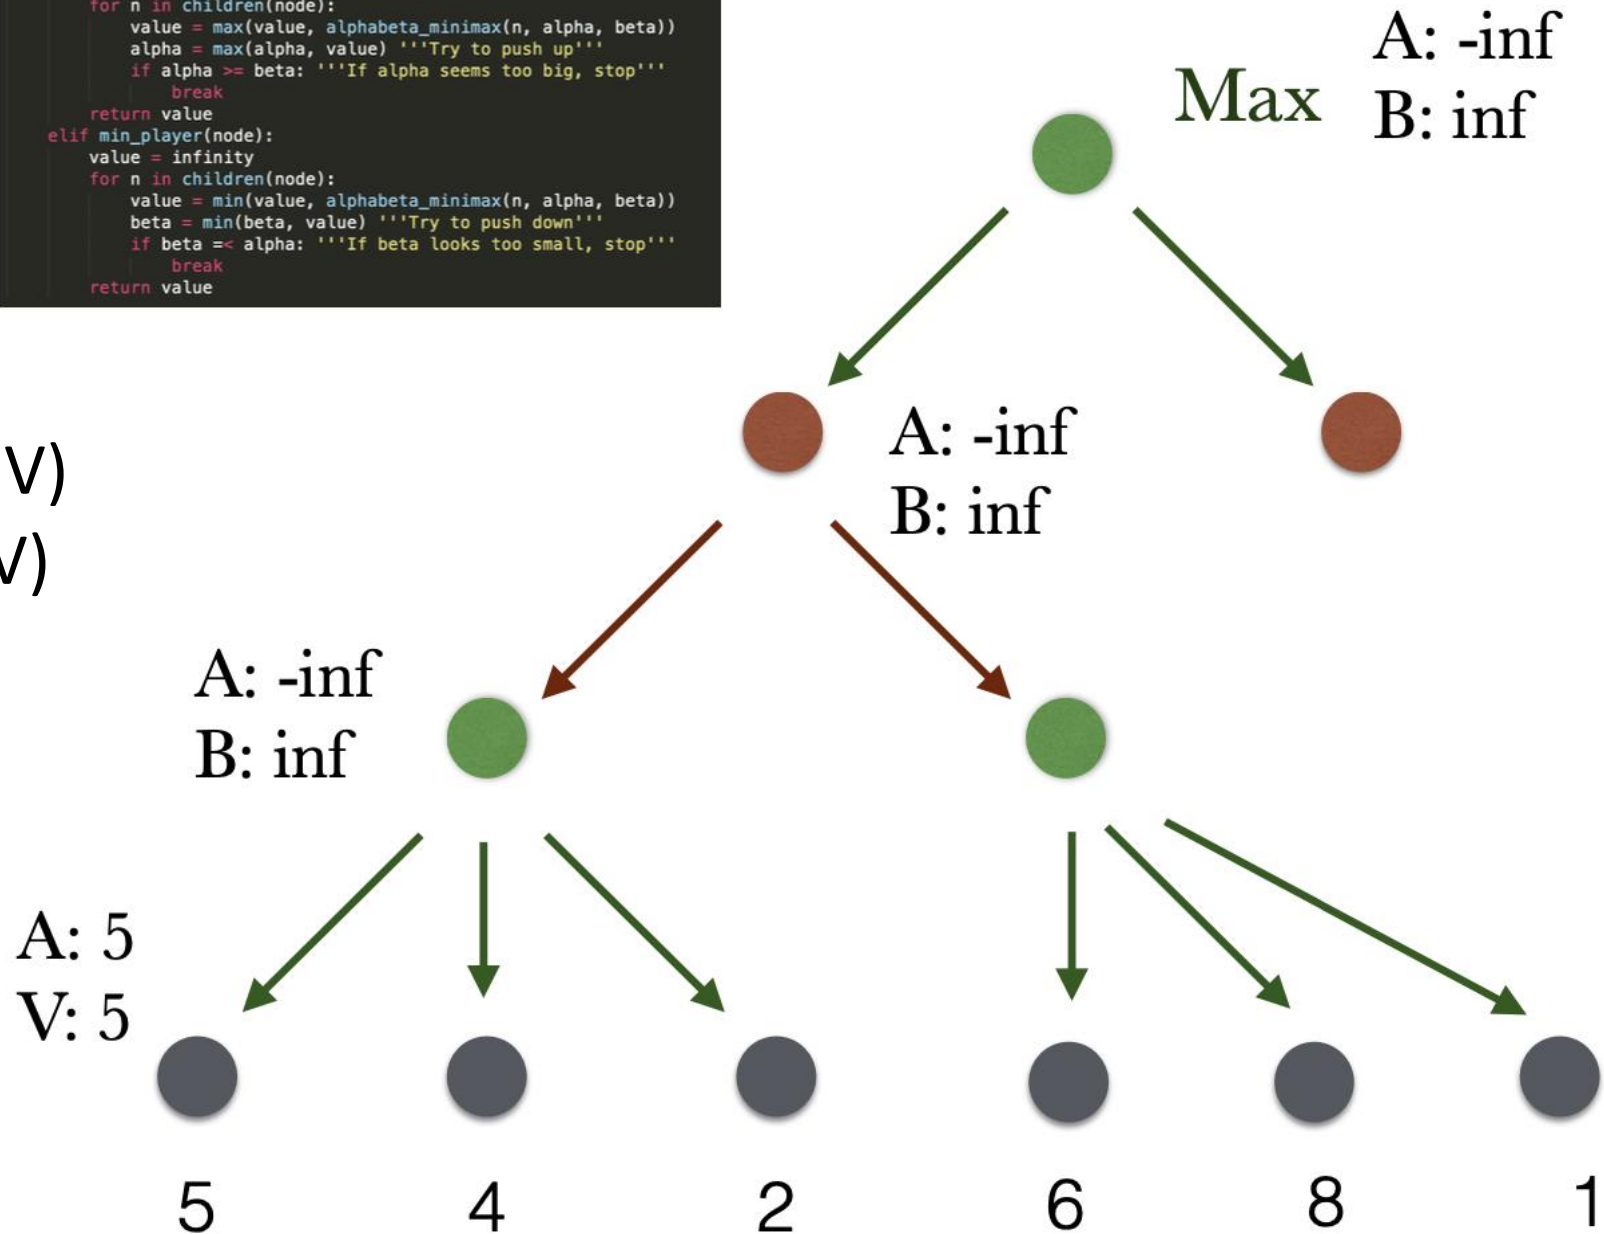


```

def alphabeta_minimax(node, alpha, beta):
    if terminal(node):
        return payoff(node)
    elif max_player(node):
        value = -infinity
        for n in children(node):
            value = max(value, alphabeta_minimax(n, alpha, beta))
            alpha = max(alpha, value) '''Try to push up'''
            if alpha >= beta: '''If alpha seems too big, stop'''
                break
        return value
    elif min_player(node):
        value = infinity
        for n in children(node):
            value = min(value, alphabeta_minimax(n, alpha, beta))
            beta = min(beta, value) '''Try to push down'''
            if beta <= alpha: '''If beta looks too small, stop'''
                break
        return value

```

$A = \max(A, V)$
 $B = \min(B, V)$
 $V = \max(V')$

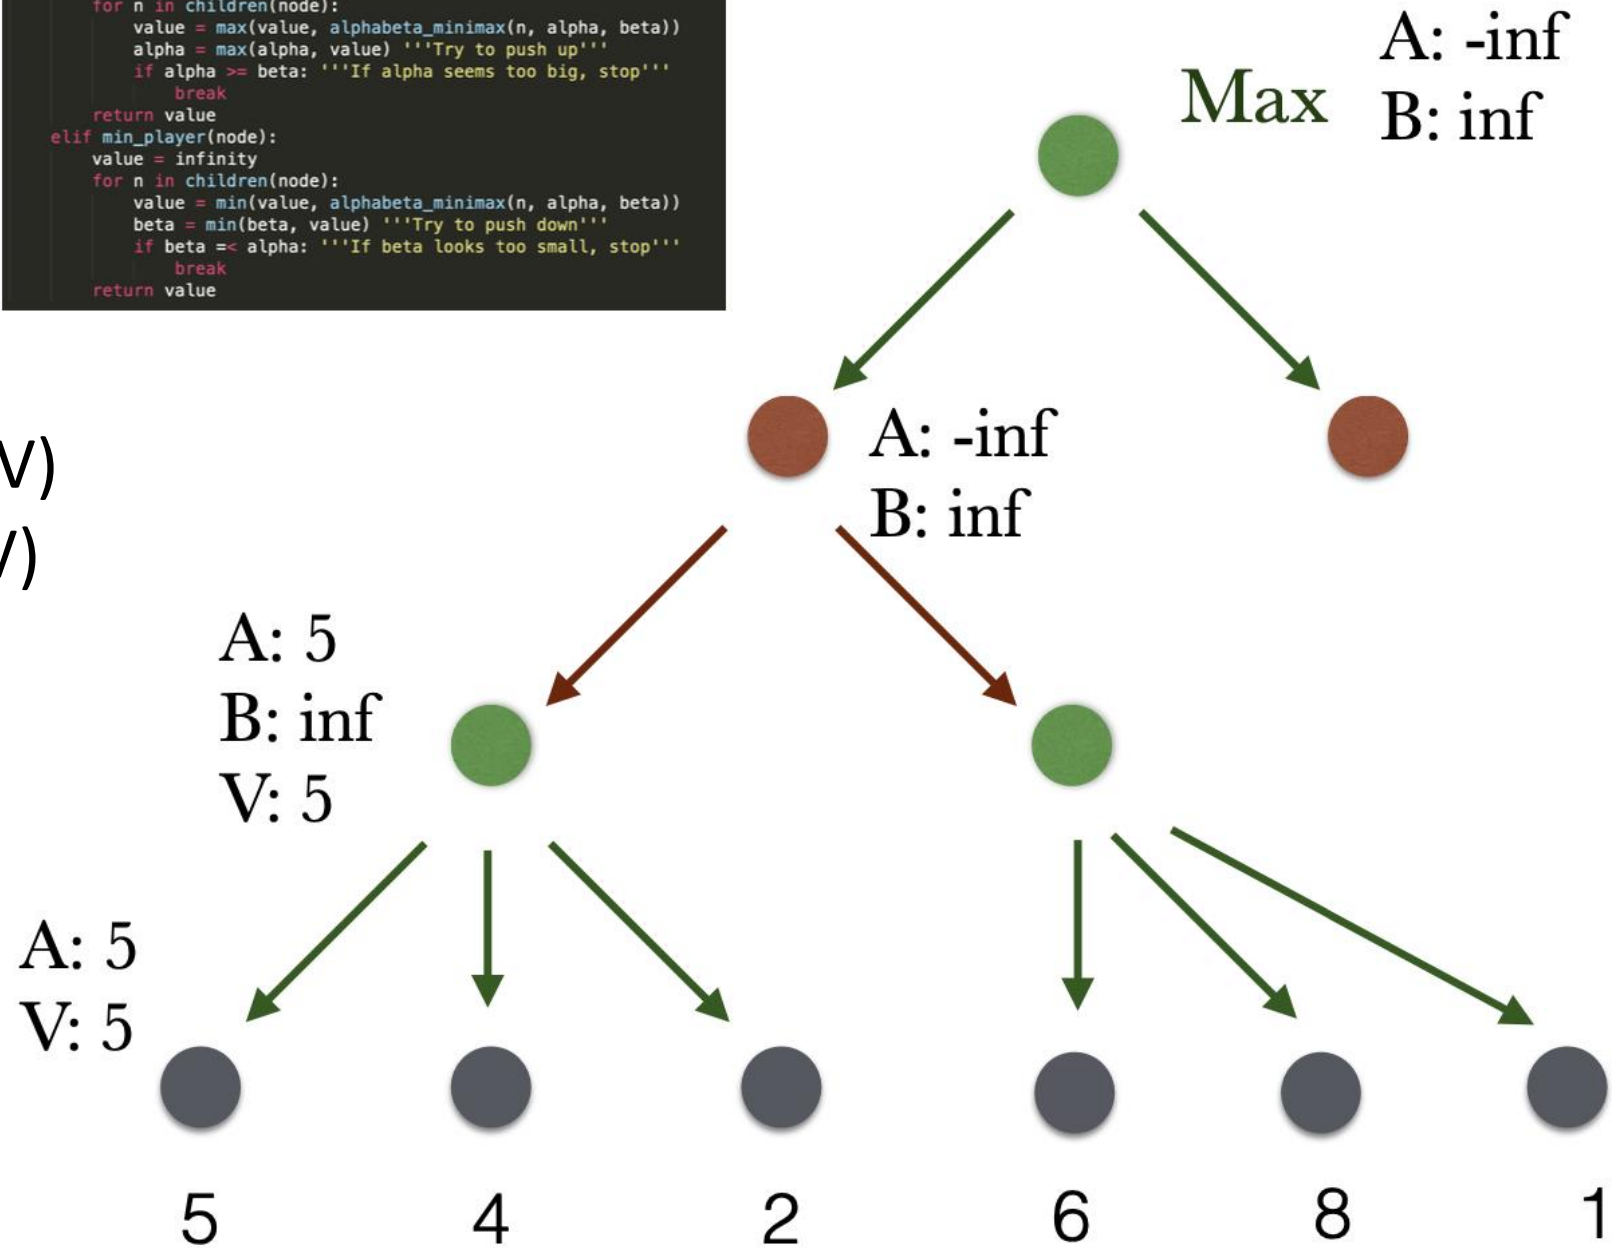


```

def alphabeta_minimax(node, alpha, beta):
    if terminal(node):
        return payoff(node)
    elif max_player(node):
        value = -infinity
        for n in children(node):
            value = max(value, alphabeta_minimax(n, alpha, beta))
            alpha = max(alpha, value)
            if alpha >= beta:
                break
        return value
    elif min_player(node):
        value = infinity
        for n in children(node):
            value = min(value, alphabeta_minimax(n, alpha, beta))
            beta = min(beta, value)
            if beta <= alpha:
                break
        return value

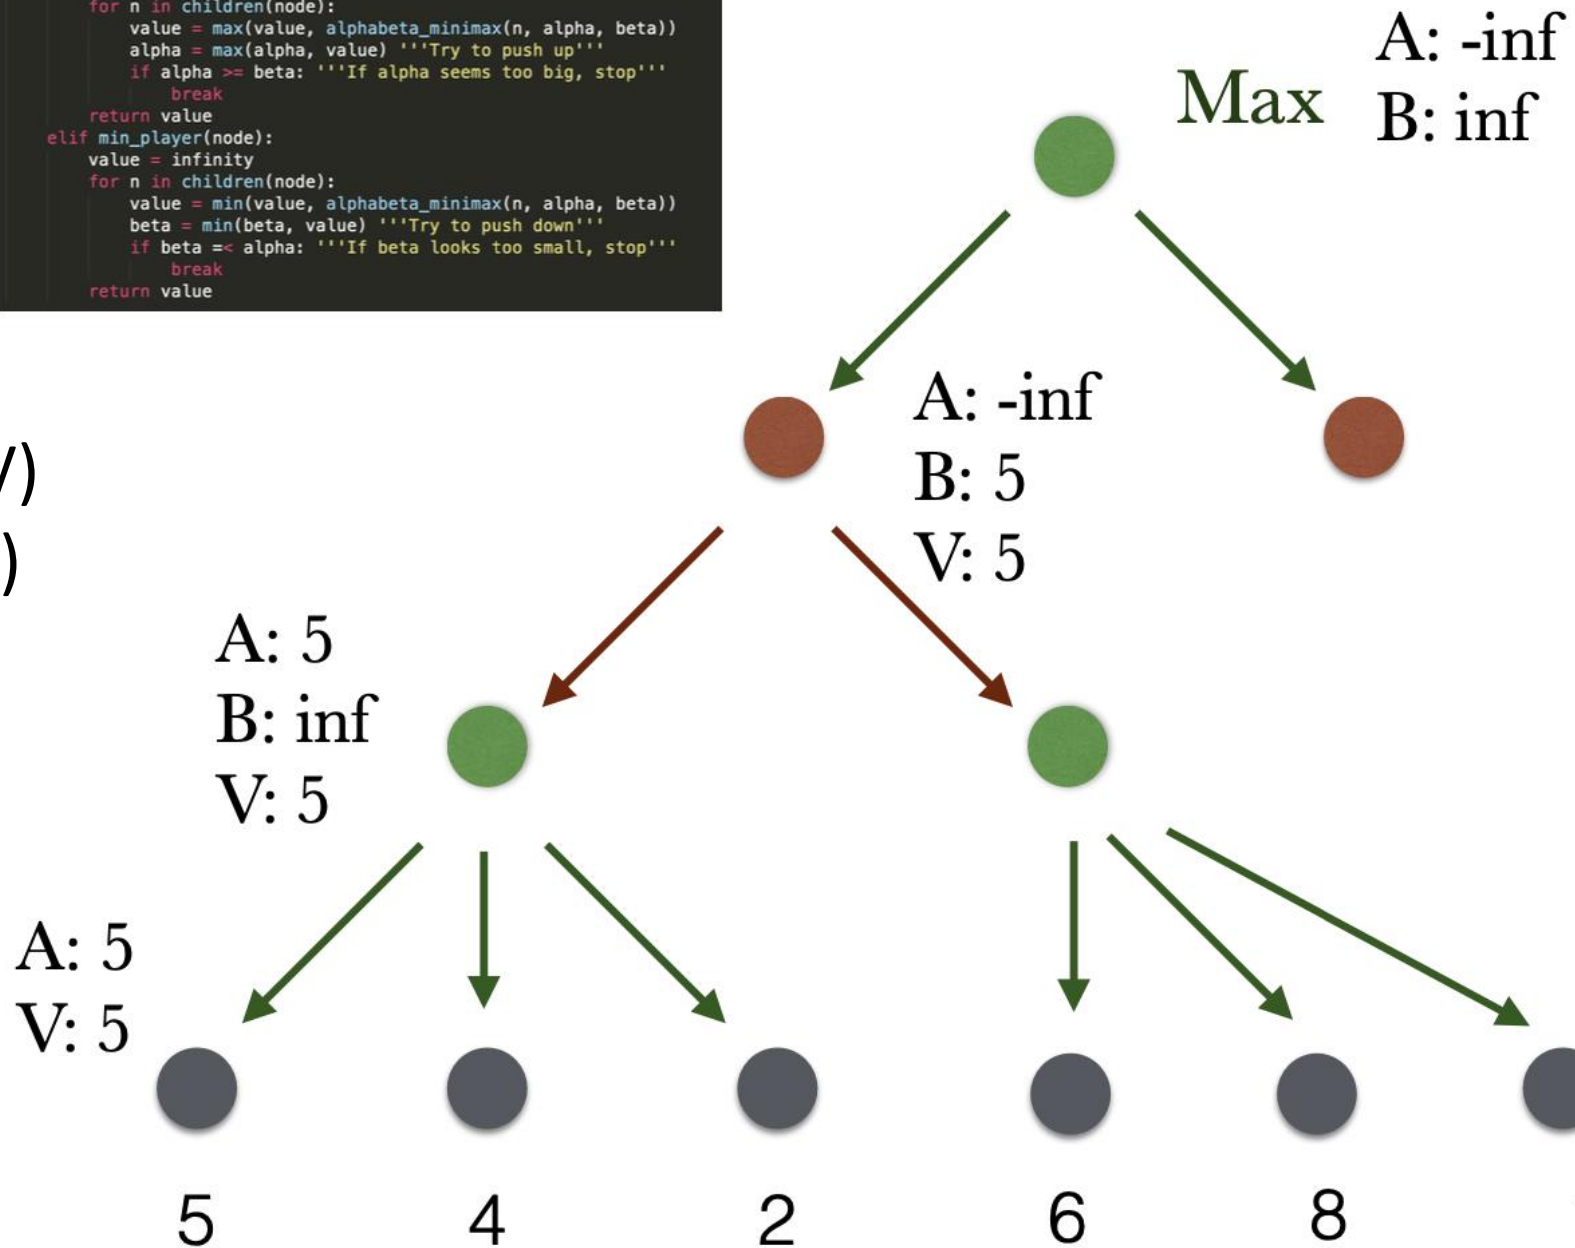
```

$A = \max(A, V)$
 $B = \min(B, V)$
 $V = \max(V')$



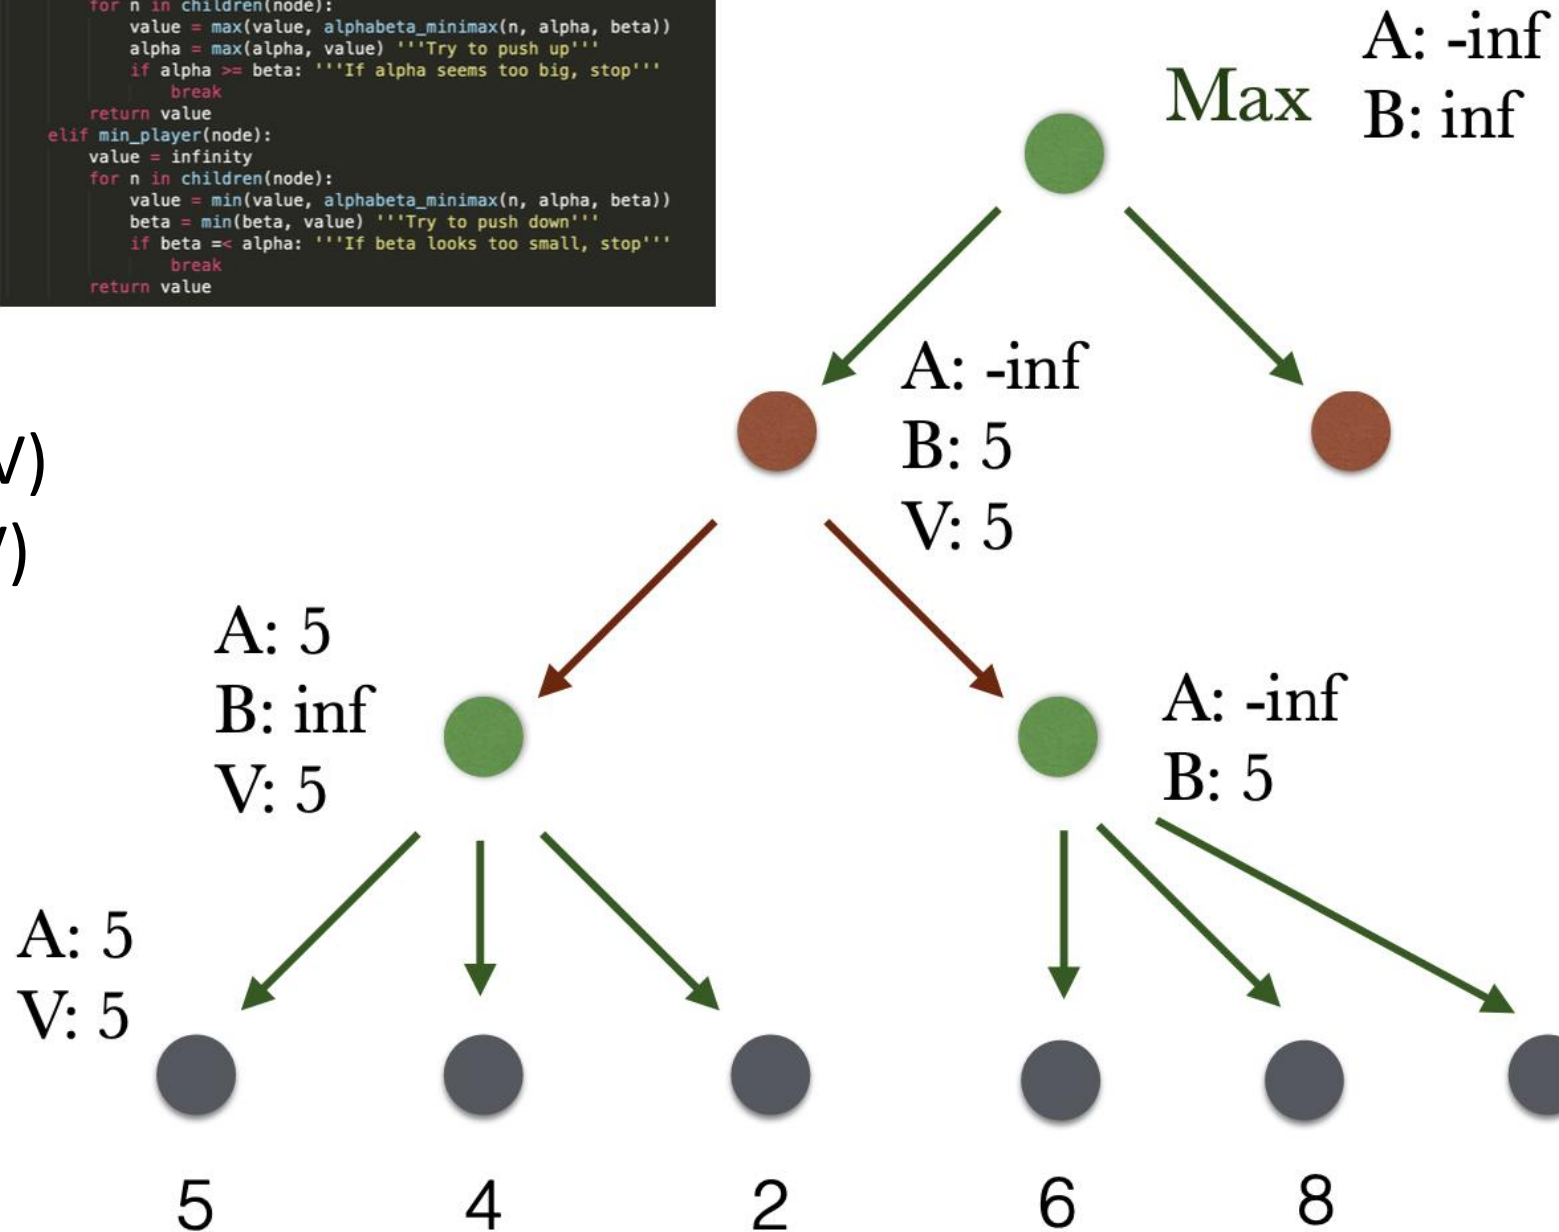
```
def alphabeta_minimax(node, alpha, beta):
    if terminal(node):
        return payoff(node)
    elif max_player(node):
        value = -infinity
        for n in children(node):
            value = max(value, alphabeta_minimax(n, alpha, beta))
            alpha = max(alpha, value) '''Try to push up'''
            if alpha >= beta: '''If alpha seems too big, stop'''
                break
        return value
    elif min_player(node):
        value = infinity
        for n in children(node):
            value = min(value, alphabeta_minimax(n, alpha, beta))
            beta = min(beta, value) '''Try to push down'''
            if beta <= alpha: '''If beta looks too small, stop'''
                break
        return value
```

$A = \max(A, V)$
 $B = \min(B, V)$
 $V = \max(V')$



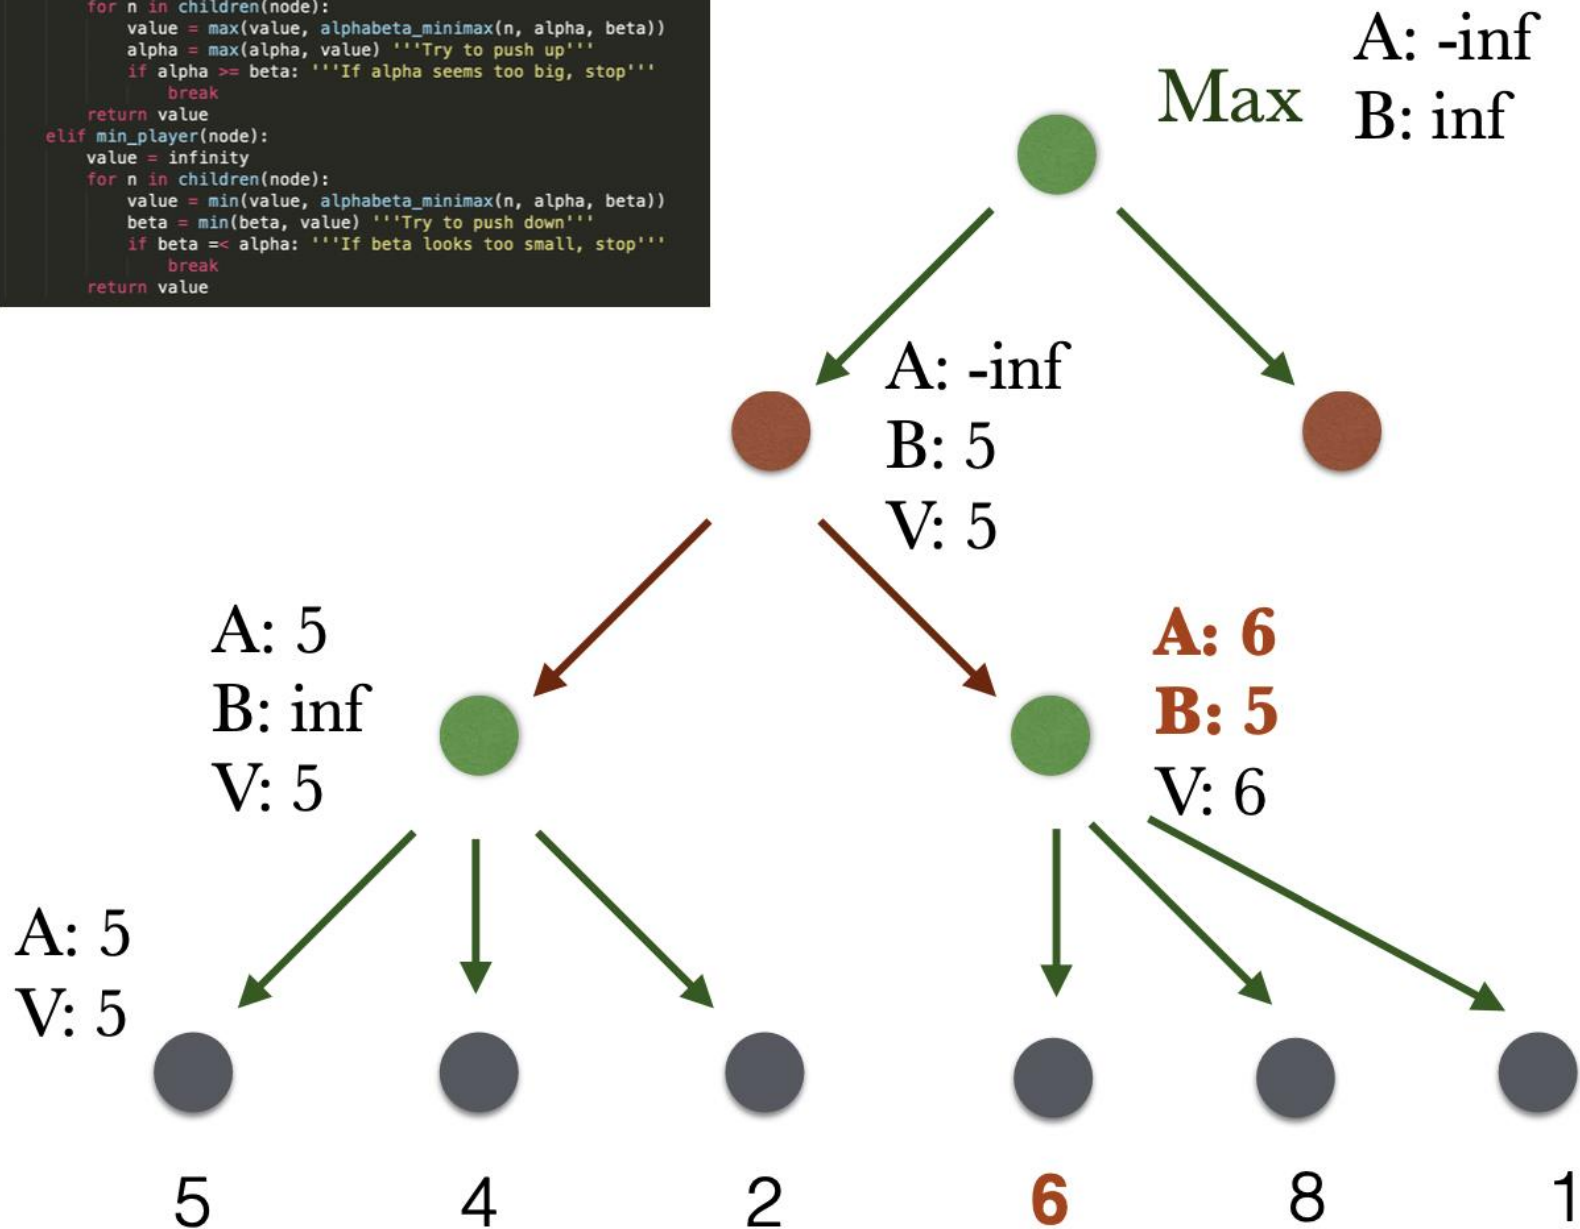
```
def alphabeta_minimax(node, alpha, beta):
    if terminal(node):
        return payoff(node)
    elif max_player(node):
        value = -infinity
        for n in children(node):
            value = max(value, alphabeta_minimax(n, alpha, beta))
            alpha = max(alpha, value) '''Try to push up'''
            if alpha >= beta: '''If alpha seems too big, stop'''
                break
        return value
    elif min_player(node):
        value = infinity
        for n in children(node):
            value = min(value, alphabeta_minimax(n, alpha, beta))
            beta = min(beta, value) '''Try to push down'''
            if beta <= alpha: '''If beta looks too small, stop'''
                break
        return value
```

$A = \max(A, V)$
 $B = \min(B, V)$
 $V = \max(V')$



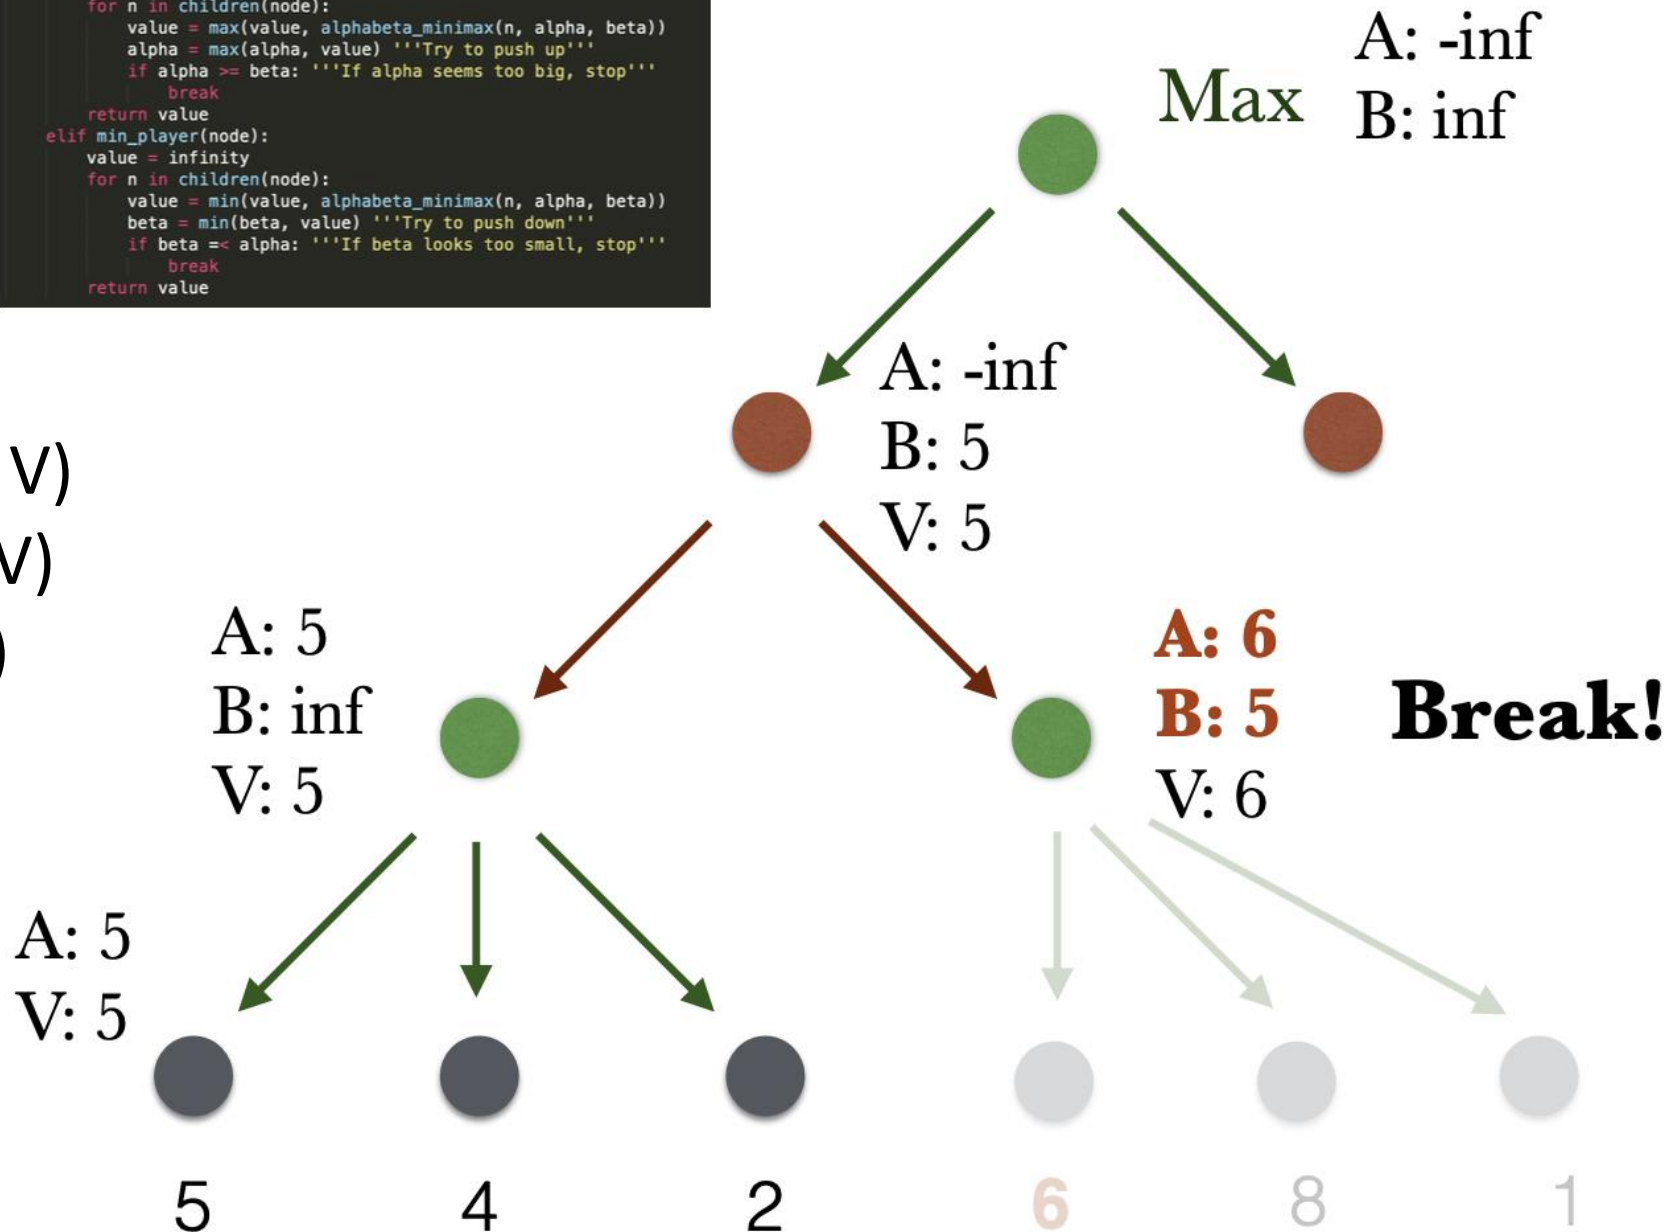
```
def alphabeta_minimax(node, alpha, beta):
    if terminal(node):
        return payoff(node)
    elif max_player(node):
        value = -infinity
        for n in children(node):
            value = max(value, alphabeta_minimax(n, alpha, beta))
            alpha = max(alpha, value) '''Try to push up'''
            if alpha >= beta: '''If alpha seems too big, stop'''
                break
        return value
    elif min_player(node):
        value = infinity
        for n in children(node):
            value = min(value, alphabeta_minimax(n, alpha, beta))
            beta = min(beta, value) '''Try to push down'''
            if beta <= alpha: '''If beta looks too small, stop'''
                break
        return value
```

$A = \max(A, V)$
 $B = \min(B, V)$
 $V = \max(V')$



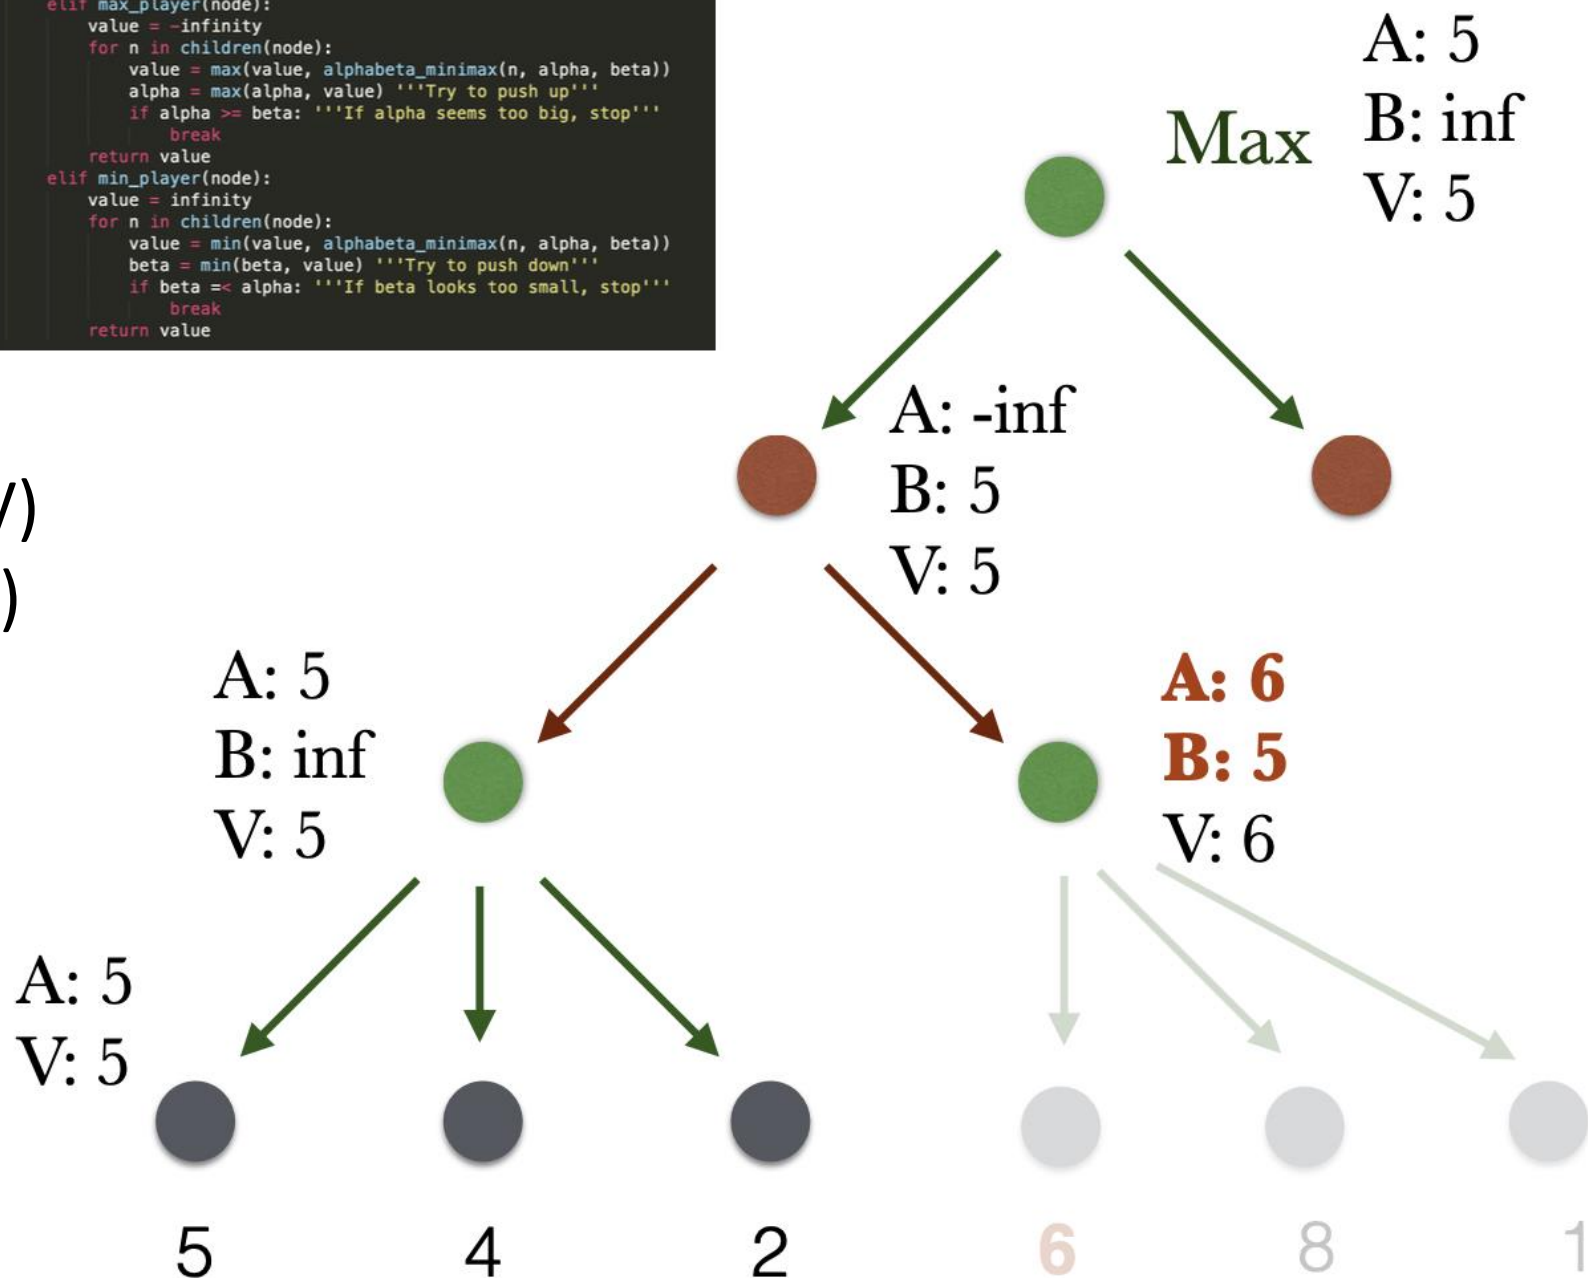
```
def alphabeta_minimax(node, alpha, beta):
    if terminal(node):
        return payoff(node)
    elif max_player(node):
        value = -infinity
        for n in children(node):
            value = max(value, alphabeta_minimax(n, alpha, beta))
            alpha = max(alpha, value)
            if alpha >= beta:
                break
        return value
    elif min_player(node):
        value = infinity
        for n in children(node):
            value = min(value, alphabeta_minimax(n, alpha, beta))
            beta = min(beta, value)
            if beta <= alpha:
                break
        return value
```

$A = \max(A, V)$
 $B = \min(B, V)$
 $V = \max(V')$



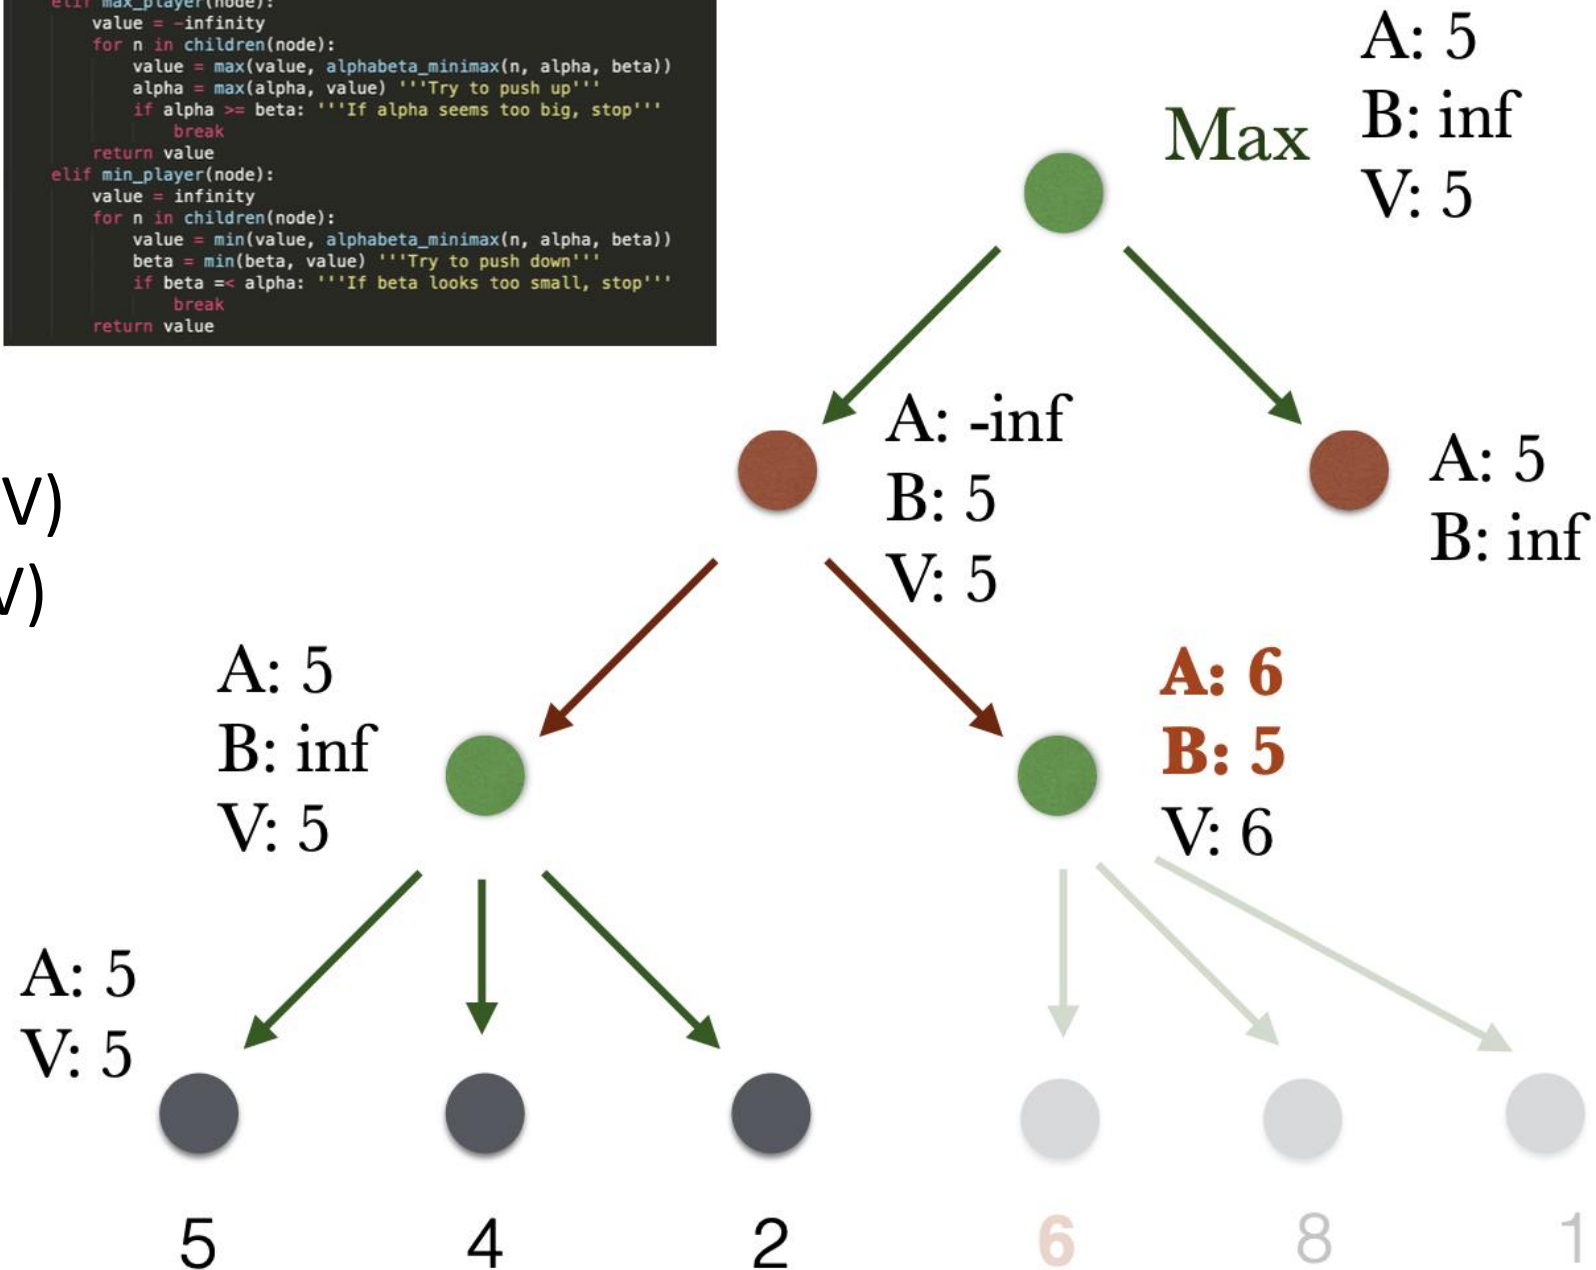

```
def alphabeta_minimax(node, alpha, beta):
    if terminal(node):
        return payoff(node)
    elif max_player(node):
        value = -infinity
        for n in children(node):
            value = max(value, alphabeta_minimax(n, alpha, beta))
            alpha = max(alpha, value)
            if alpha >= beta:
                break
        return value
    elif min_player(node):
        value = infinity
        for n in children(node):
            value = min(value, alphabeta_minimax(n, alpha, beta))
            beta = min(beta, value)
            if beta <= alpha:
                break
        return value
```

$A = \max(A, V)$
 $B = \min(B, V)$
 $V = \max(V')$




```
def alphabeta_minimax(node, alpha, beta):
    if terminal(node):
        return payoff(node)
    elif max_player(node):
        value = -infinity
        for n in children(node):
            value = max(value, alphabeta_minimax(n, alpha, beta))
            alpha = max(alpha, value)
            if alpha >= beta:
                break
        return value
    elif min_player(node):
        value = infinity
        for n in children(node):
            value = min(value, alphabeta_minimax(n, alpha, beta))
            beta = min(beta, value)
            if beta <= alpha:
                break
        return value
```


$A = \max(A, V)$
 $B = \min(B, V)$
 $V = \max(V')$



Vertical Cut-off

Cut-off at some fixed depth and use heuristics to evaluate the intermediate nodes.





Evaluation function: score + more heuristics

Expectimax

- Instead of MIN PLAYER, we have **CHANCE PLAYER**
- At each CHANCE PLAYER's state, the action to take is chosen from the set of all possible actions **with probabilities**

Expectimax

- Instead of MIN PLAYER, we have **CHANCE PLAYER**
- At each CHANCE PLAYER's state, the action to take is chosen from the set of all possible actions **with probabilities**

~~Min Node: $\min(s_1^t, s_2^t, \dots, s_n^t)$~~ → Chance Node: $\sum_{s'} P(s') \cdot V(s')$ (Expected Value)

PA: 2048


- MAX PLAYER: 4 possible moves → **{Up, Down, Left, Right}**
(Some moves might be invalid → Remove them from the tree)
- CHANCE PLAYER (Agent): Choose the next state **uniformly random** from the empty spots

PA: 2048 (build_tree)

- MAX PLAYER: 4 possible moves
→ {Up, Down, Left, Right}

(Some moves might be invalid
→ Remove them from the tree)

- CHANCE PLAYER (Agent):
Choose the next state
uniformly random from the
empty spots




(Max Player)

PA: 2048 (build_tree)

- MAX PLAYER: 4 possible moves
→ {Up, Down, Left, Right}

(Some moves might be invalid
→ Remove them from the tree)

- CHANCE PLAYER (Agent):
Choose the next state
uniformly random from the
empty spots




PA: 2048 (build_tree)

- MAX PLAYER: 4 possible moves
→ {Up, Down, Left, Right}

(Some moves might be invalid
→ Remove them from the tree)

- CHANCE PLAYER (Agent):
Choose the next state
uniformly random from the
empty spots




PA: 2048 (build_tree)

- MAX PLAYER: 4 possible moves
→ {Up, Down, Left, Right}

(Some moves might be invalid
→ Remove them from the tree)

- CHANCE PLAYER (Agent):
Choose the next state
uniformly random from the
empty spots

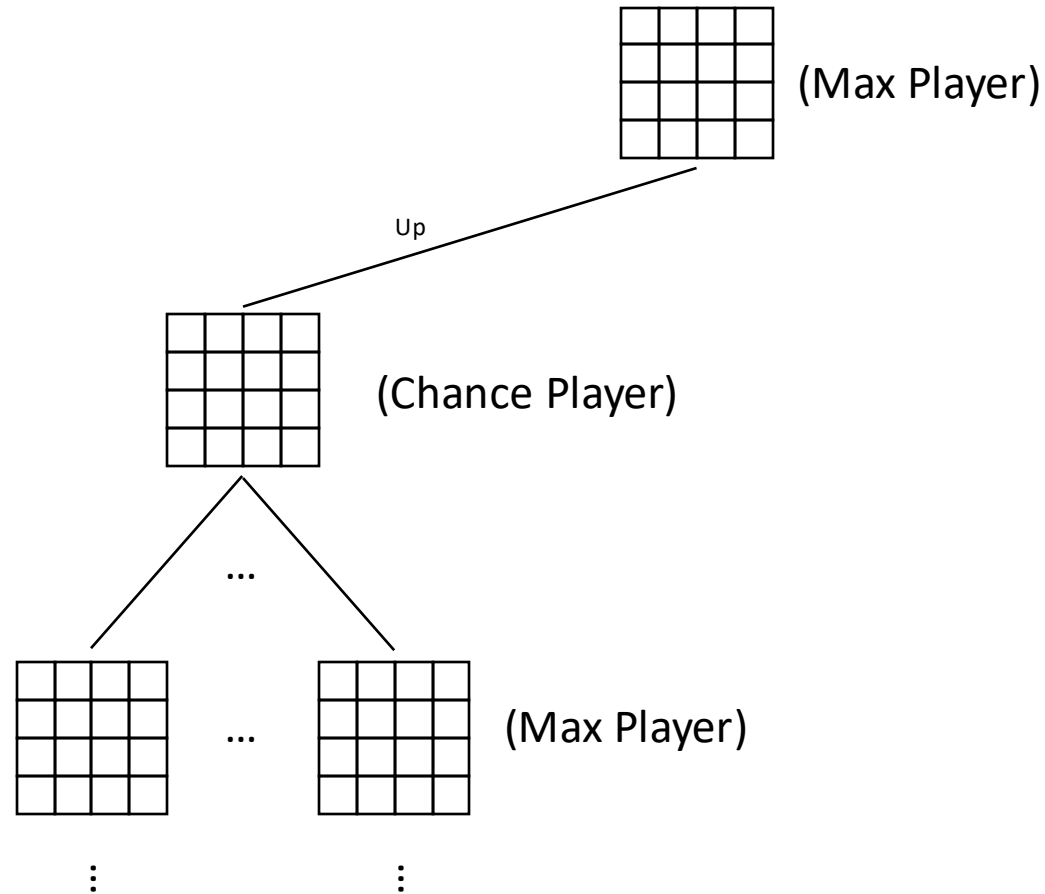


PA: 2048 (build_tree)

- MAX PLAYER: 4 possible moves
→ {Up, Down, Left, Right}

(Some moves might be invalid
→ Remove them from the tree)

- CHANCE PLAYER (Agent):
Choose the next state
uniformly random from the
empty spots

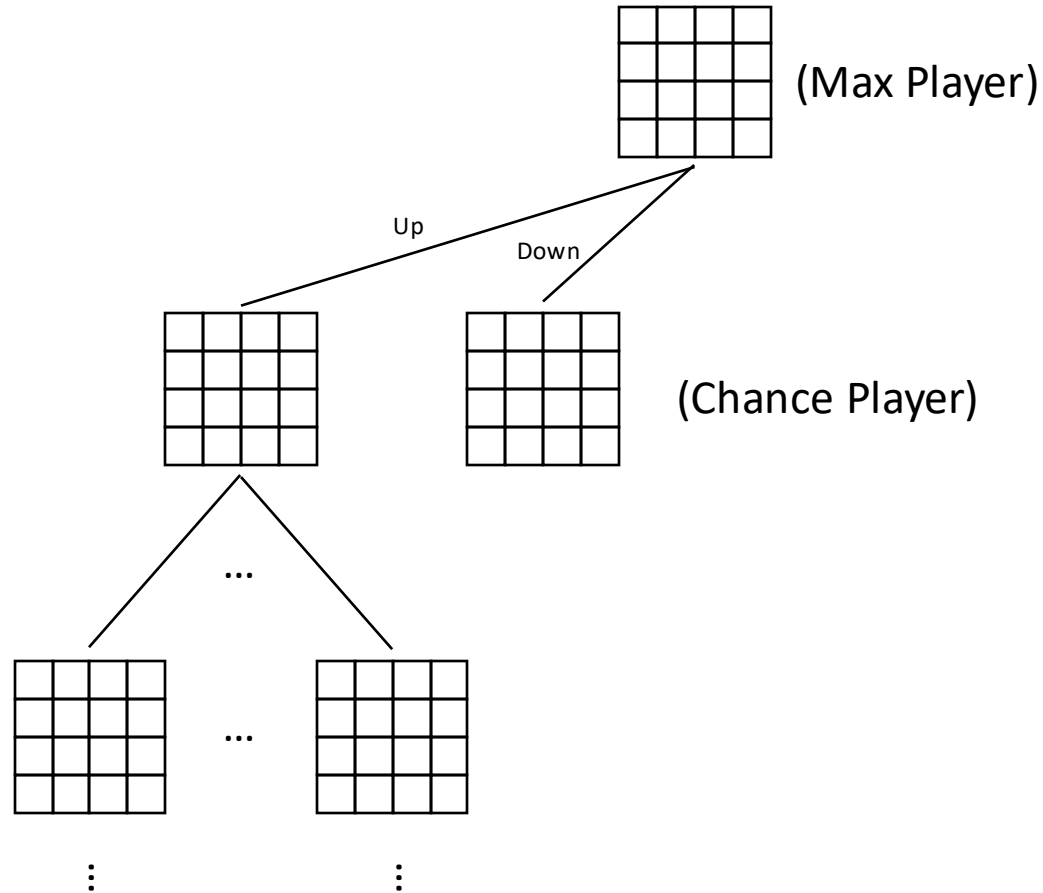


PA: 2048 (build_tree)

- MAX PLAYER: 4 possible moves
→ {Up, Down, Left, Right}

(Some moves might be invalid
→ Remove them from the tree)

- CHANCE PLAYER (Agent):
Choose the next state
uniformly random from the
empty spots

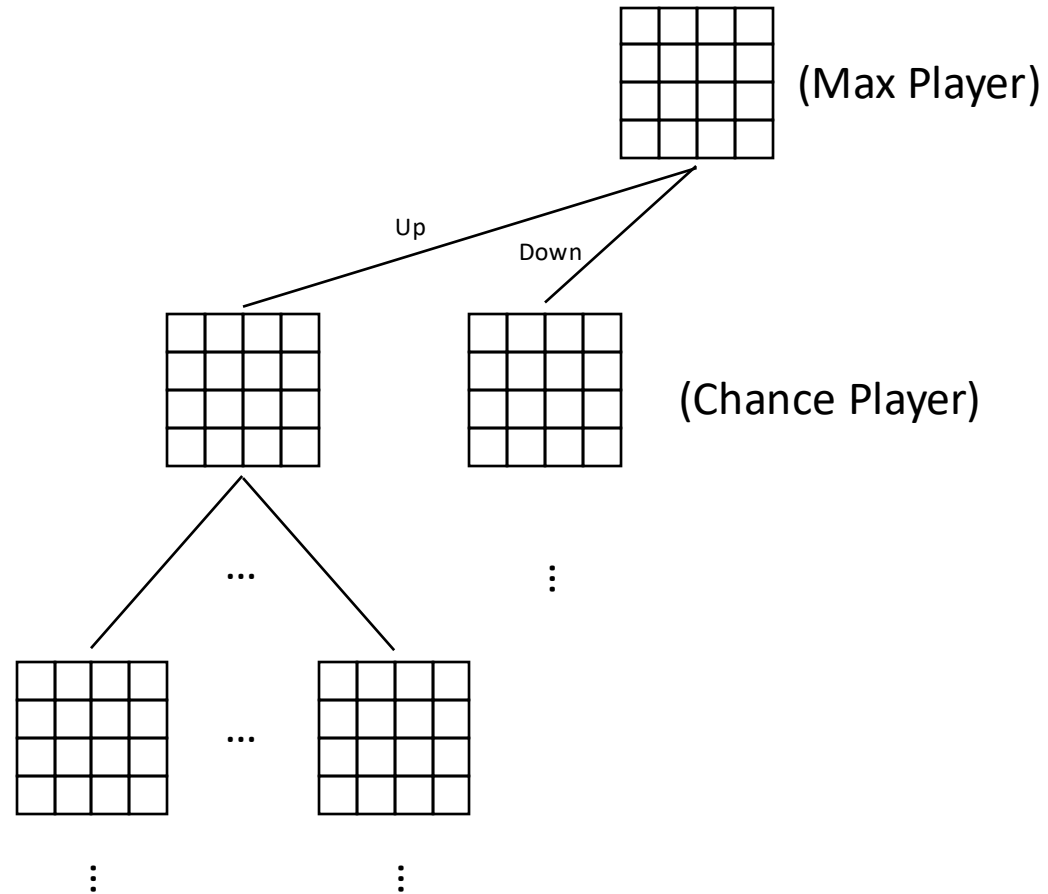


PA: 2048 (build_tree)

- MAX PLAYER: 4 possible moves
→ {Up, Down, Left, Right}

(Some moves might be invalid
→ Remove them from the tree)

- CHANCE PLAYER (Agent):
Choose the next state
uniformly random from the
empty spots

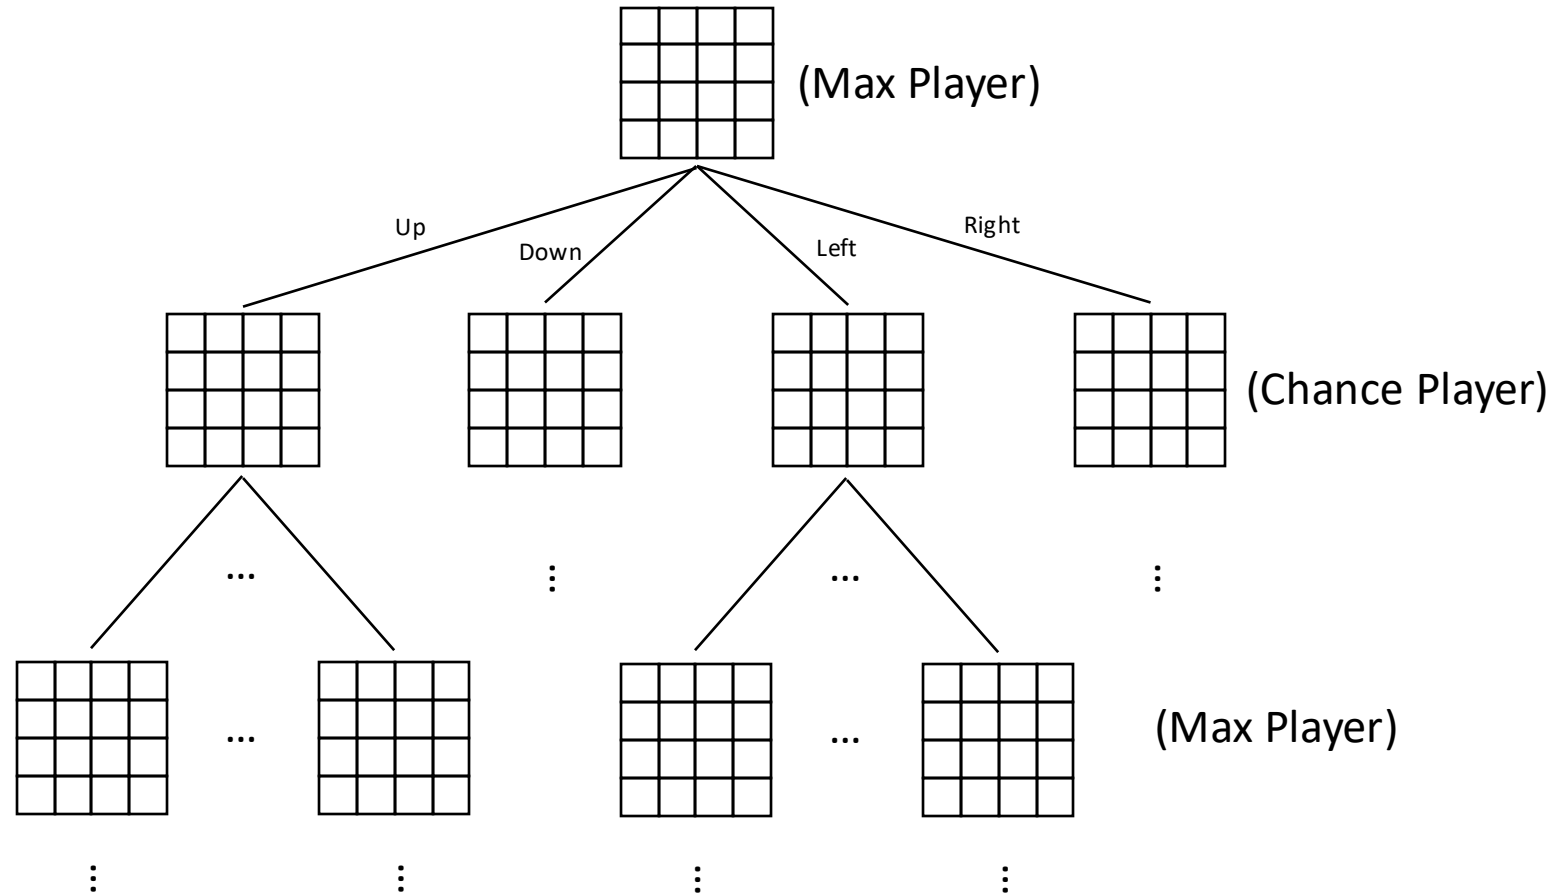


PA: 2048 (build_tree)

- MAX PLAYER: 4 possible moves
→ {Up, Down, Left, Right}

(Some moves might be invalid
→ Remove them from the tree)

- CHANCE PLAYER (Agent):
Choose the next state
uniformly random from the
empty spots

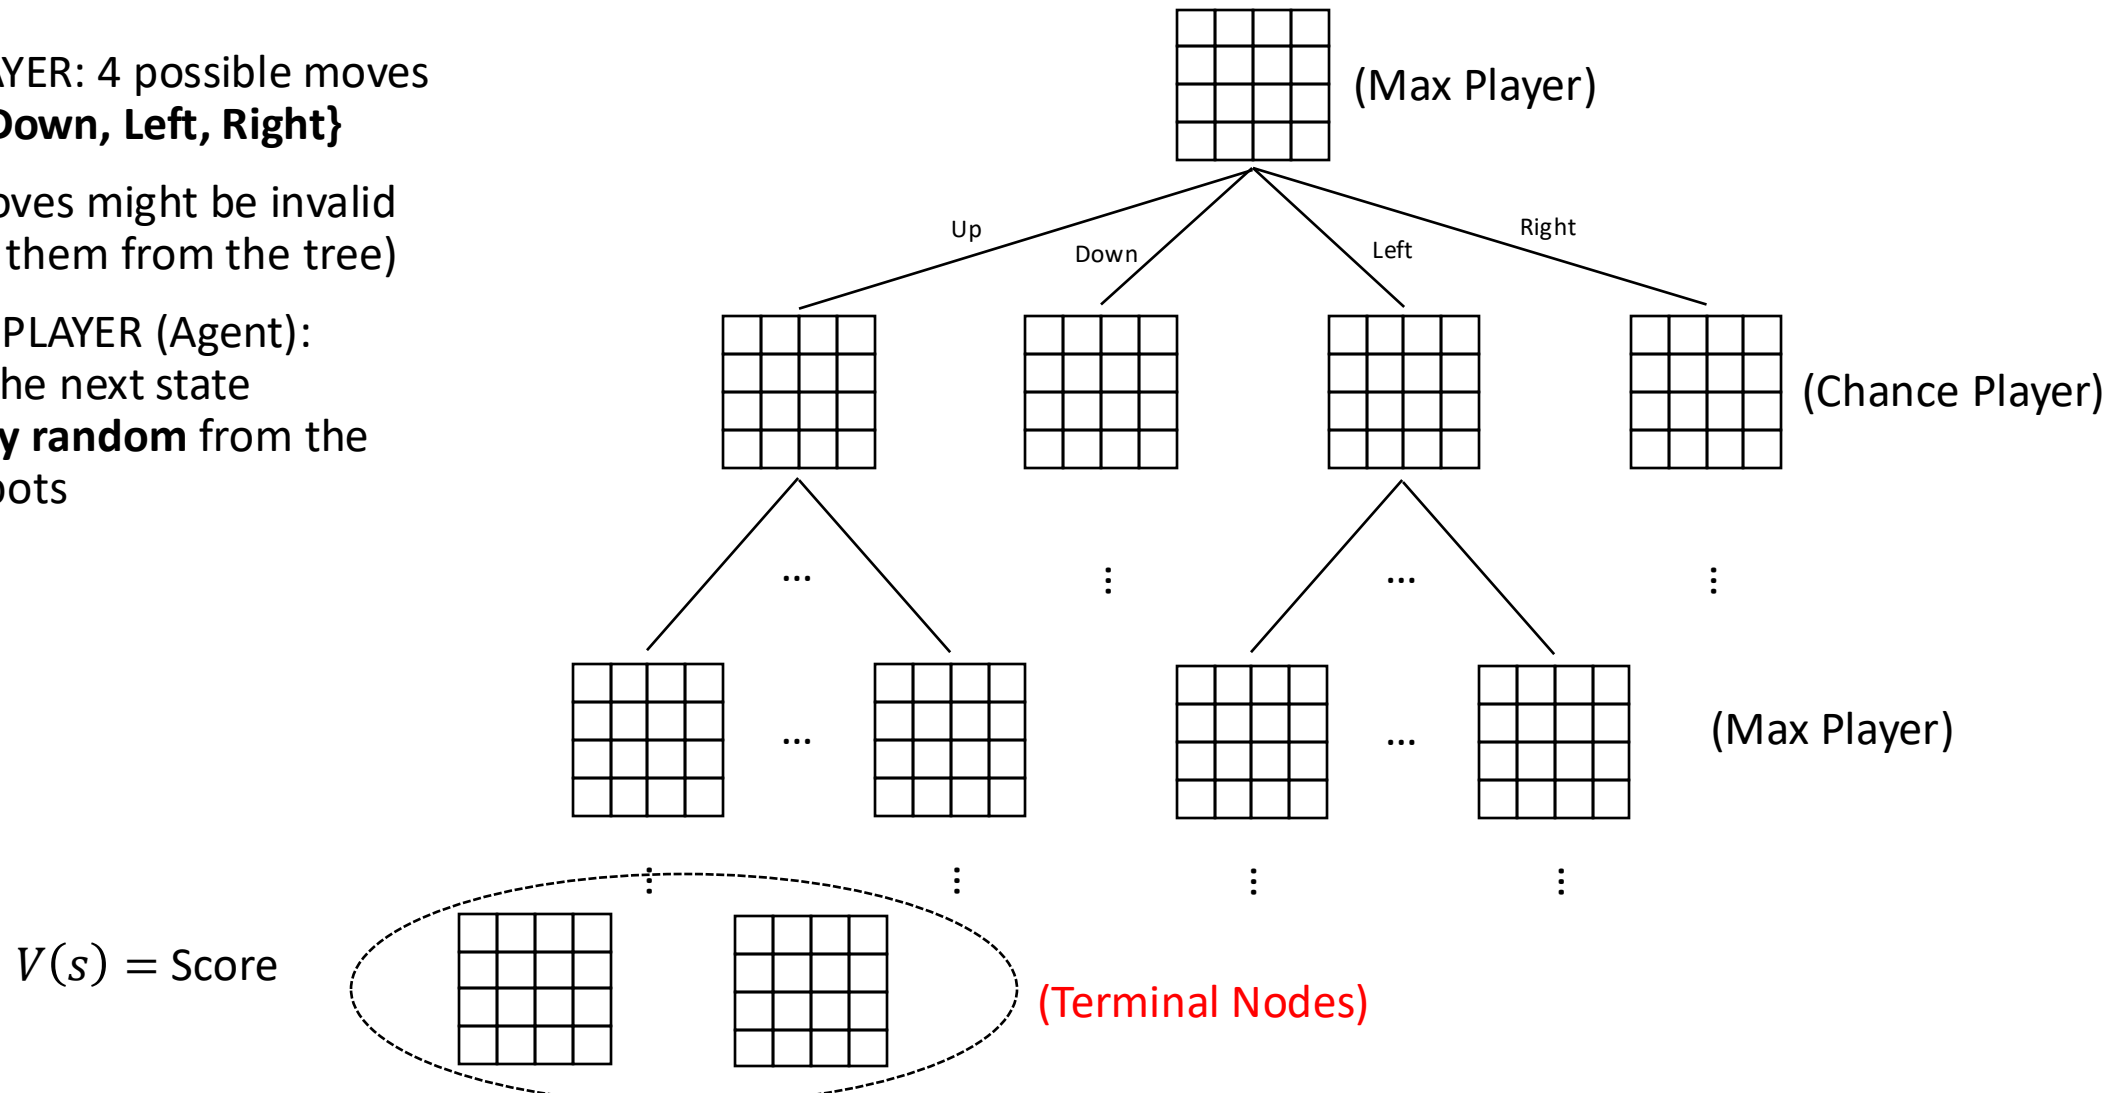


PA: 2048 (expectimax)

- MAX PLAYER: 4 possible moves
→ {Up, Down, Left, Right}

(Some moves might be invalid
→ Remove them from the tree)

- CHANCE PLAYER (Agent):
Choose the next state
uniformly random from the
empty spots




PA: 2048 (expectimax)

- MAX PLAYER: 4 possible moves
→ {Up, Down, Left, Right}

(Some moves might be invalid
→ Remove them from the tree)

- CHANCE PLAYER (Agent):
Choose the next state
uniformly random from the
empty spots




PA: 2048 (expectimax)

- MAX PLAYER: 4 possible moves
→ {Up, Down, Left, Right}

(Some moves might be invalid
→ Remove them from the tree)

- CHANCE PLAYER (Agent):
Choose the next state
uniformly random from the
empty spots



PA: 2048

Helpful Functions in `game.py`:

- `current_state()`: Returns the current state (`tile_matrix`, `score`)
- `move(direction)`: Returns `True` if the move was successfully taken
- `set_state(tile_matrix, score)`: Sets the state of the game
- `get_open_tiles()`: Returns a list of empty positions
[(`__`, `__`), (`__`, `__`), ..., (`__`, `__`)]
- `Game(tile_matrix, score)`: Creates a new game object (useful for “simulating” moves)

PA: 2048

| Score: 0 | | | |
|----------|---|---|---|
| 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 |
| 2 | 0 | 0 | 0 |
| 0 | 0 | 0 | 2 |

PA: 2048 Extra Credit

- Good luck 😊
- Teaching staff will not be giving hints ☹️